

statistical-analysis

September 8, 2024

```
[4]: import pandas as pd

# Load the dataset
df = pd.read_csv("C:/Users/user/OneDrive/Documents/sleep_health_dataset.csv")

# Show the first few rows to confirm it loaded correctly
df.head()
```

```
[4]:
```

	Person ID	Age	Sleep Duration	Quality of Sleep	Physical Activity Level	\
0	1	27	6.1	6	42	
1	2	28	6.2	6	60	
2	3	28	6.2	6	60	
3	4	28	5.9	4	30	
4	5	28	5.9	4	30	

	Stress Level	High BP	Low BP	Daily Steps
0	6	126	83	4200
1	8	125	80	10000
2	8	125	80	10000
3	8	140	90	3000
4	8	140	90	3000

```
[5]: # Calculate descriptive statistics for numerical columns
relevant_features = ['Age', 'Sleep Duration', 'Quality of Sleep', 'Physical_
↳Activity Level',
                    'Stress Level', 'High BP', 'Low BP', 'Daily Steps']

# Mean
mean_values = df[relevant_features].mean()

# Median
median_values = df[relevant_features].median()

# Mode
mode_values = df[relevant_features].mode().iloc[0] # mode() returns a
↳dataframe, selecting the first mode
```

```

# Standard Deviation
std_values = df[relevant_features].std()

# Variance
variance_values = df[relevant_features].var()

# Print the results
print("Mean Values:\n", mean_values)
print("\nMedian Values:\n", median_values)
print("\nMode Values:\n", mode_values)
print("\nStandard Deviation:\n", std_values)
print("\nVariance:\n", variance_values)

```

Mean Values:

Age	42.184492
Sleep Duration	7.132086
Quality of Sleep	7.312834
Physical Activity Level	59.171123
Stress Level	5.385027
High BP	128.553476
Low BP	84.649733
Daily Steps	6816.844920

dtype: float64

Median Values:

Age	43.0
Sleep Duration	7.2
Quality of Sleep	7.0
Physical Activity Level	60.0
Stress Level	5.0
High BP	130.0
Low BP	85.0
Daily Steps	7000.0

dtype: float64

Mode Values:

Age	43.0
Sleep Duration	7.2
Quality of Sleep	8.0
Physical Activity Level	60.0
Stress Level	3.0
High BP	130.0
Low BP	80.0
Daily Steps	8000.0

Name: 0, dtype: float64

Standard Deviation:

```

Age                8.673133
Sleep Duration     0.795657
Quality of Sleep   1.196956
Physical Activity Level 20.830804
Stress Level       1.774526
High BP            7.748118
Low BP             6.161611
Daily Steps        1617.915679
dtype: float64

```

```

Variance:
Age                7.522324e+01
Sleep Duration     6.330696e-01
Quality of Sleep   1.432703e+00
Physical Activity Level 4.339224e+02
Stress Level       3.148944e+00
High BP            6.003333e+01
Low BP             3.796546e+01
Daily Steps        2.617651e+06
dtype: float64

```

```

[6]: from scipy import stats

# Null hypothesis: mean sleep duration = 7 hours
# Alternative hypothesis: mean sleep duration != 7 hours
sleep_duration = df['Sleep Duration']
t_stat, p_value = stats.ttest_1samp(sleep_duration, 7)

print(f"T-statistic: {t_stat}, P-value: {p_value}")

# Interpret the result
alpha = 0.05 # significance level
if p_value < alpha:
    print("Reject the null hypothesis - The mean sleep duration is_
    ↪significantly different from 7 hours.")
else:
    print("Fail to reject the null hypothesis - The mean sleep duration is not_
    ↪significantly different from 7 hours.")

```

```

T-statistic: 3.2104462758942, P-value: 0.0014402421900475528
Reject the null hypothesis - The mean sleep duration is significantly different
from 7 hours.

```

```

[7]: import numpy as np

# Function to compute confidence interval
def confidence_interval(data, confidence=0.95):

```

```

mean = np.mean(data)
std_err = stats.sem(data) # standard error
margin_of_error = std_err * stats.t.ppf((1 + confidence) / 2, len(data) - 1)
return mean - margin_of_error, mean + margin_of_error

# Compute 95% confidence interval for daily steps
ci_lower, ci_upper = confidence_interval(df['Daily Steps'])

print(f"95% Confidence Interval for Daily Steps: ({ci_lower}, {ci_upper})")

```

95% Confidence Interval for Daily Steps: (6652.339714058175, 6981.350125514018)

```

[8]: import statsmodels.api as sm

# Define independent and dependent variables
X = df['Age'] # Independent variable
y = df['High BP'] # Dependent variable

# Add a constant to the independent variable (intercept)
X = sm.add_constant(X)

# Fit the regression model
model = sm.OLS(y, X).fit()

# Output the summary of the regression
print(model.summary())

```

```

OLS Regression Results
=====
Dep. Variable:          High BP    R-squared:                0.367
Model:                  OLS        Adj. R-squared:           0.365
Method:                 Least Squares    F-statistic:             215.8
Date:                  Sun, 08 Sep 2024    Prob (F-statistic):       7.62e-39
Time:                  21:32:55    Log-Likelihood:          -1210.4
No. Observations:      374    AIC:                     2425.
Df Residuals:          372    BIC:                     2433.
Df Model:              1
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	105.7207	1.587	66.622	0.000	102.600	108.841
Age	0.5413	0.037	14.689	0.000	0.469	0.614

```

=====
Omnibus:                7.988    Durbin-Watson:           0.871
Prob(Omnibus):          0.018    Jarque-Bera (JB):        7.892
Skew:                   -0.347    Prob(JB):                0.0193
Kurtosis:               3.153    Cond. No.:               214.
=====

```

=====

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
[10]: import matplotlib.pyplot as plt
import numpy as np

# Scatter plot using matplotlib
plt.figure(figsize=(8,6))
plt.scatter(df['Age'], df['High BP'], color='blue', label='Data points')

# Fit the regression line
slope, intercept = np.polyfit(df['Age'], df['High BP'], 1) # 1 represents a linear fit

# Create the regression line
reg_line = slope * df['Age'] + intercept

# Plot the regression line
plt.plot(df['Age'], reg_line, color='red', label='Regression line')

# Add title and labels
plt.title('Relationship between Age and High Blood Pressure')
plt.xlabel('Age')
plt.ylabel('High Blood Pressure')

# Add a legend
plt.legend()

# Show the plot
plt.show()
```

